

Rain Prediction ANN

This is the model predict Wheather in next day
there will be rain or not

```
In [1]: # import all libraries
import matplotlib.pyplot as plt # for visualisation
import seaborn as sns # for statistical visualisation
import datetime # for date and time
from sklearn.preprocessing import LabelEncoder # for encode the categorical value
from sklearn.preprocessing import StandardScaler # for scale the value to -3 to 3
from sklearn.model_selection import train_test_split # to split the dependent and independent variables
# Dense: for connected layer to the neuron
# BatchNormalization: Normalizes the inputs of each layer to improve training speed
# Dropout: Prevents overfitting by randomly dropping a fraction of neurons during training
# LSTM: A layer for Long Short-Term Memory networks, designed for sequential data
from keras.layers import Dense, BatchNormalization, Dropout, LSTM
from keras.models import Sequential #Used to define a linear stack of layers.
from keras.utils import to_categorical #Converts class labels (integers) into one-hot encoding
from keras.optimizers import Adam # optimizer for ANN
from tensorflow.keras import regularizers # feature selection
from sklearn.metrics import precision_score, recall_score, confusion_matrix, classification_report
from keras import callbacks
import pandas as pd # for dataset
import numpy as np # for nD array

# remove warning
import warnings
warnings.filterwarnings("ignore")
```


```
In [2]: # Load the dataset
data=pd.read_csv(r"C:\Users\sunil\Downloads\weatherAUS.csv\weatherAUS.csv")
```

```
In [3]: data.head() # head of dataset
```

```
Out[3]:
```

	Date	Location	MinTemp	MaxTemp	Rainfall	Evaporation	Sunshine	WindGustDir
0	2008-12-01	Albury	13.4	22.9	0.6	NaN	NaN	W
1	2008-12-02	Albury	7.4	25.1	0.0	NaN	NaN	WNW
2	2008-12-03	Albury	12.9	25.7	0.0	NaN	NaN	WSW
3	2008-12-04	Albury	9.2	28.0	0.0	NaN	NaN	NE
4	2008-12-05	Albury	17.5	32.3	1.0	NaN	NaN	W

5 rows × 23 columns



The dataset contains about 10 years of daily weather observations from different locations across Australia. Observations were drawn from numerous weather stations.

In this project, I will use this data to predict whether or not it will rain the next day. There are 23 attributes including the target variable "RainTomorrow", indicating whether or not it will rain the next day or not.

```
In [4]: data.info() # information about the data
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 145460 entries, 0 to 145459
Data columns (total 23 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Date                  145460 non-null object
1   Location              145460 non-null object
2   MinTemp               143975 non-null float64
3   MaxTemp               144199 non-null float64
4   Rainfall              142199 non-null float64
5   Evaporation           82670 non-null float64
6   Sunshine              75625 non-null float64
7   WindGustDir           135134 non-null object
8   WindGustSpeed         135197 non-null float64
9   WindDir9am            134894 non-null object
10  WindDir3pm            141232 non-null object
11  WindSpeed9am          143693 non-null float64
12  WindSpeed3pm          142398 non-null float64
13  Humidity9am           142806 non-null float64
14  Humidity3pm           140953 non-null float64
15  Pressure9am           130395 non-null float64
16  Pressure3pm           130432 non-null float64
17  Cloud9am              89572 non-null float64
18  Cloud3pm              86102 non-null float64
19  Temp9am               143693 non-null float64
20  Temp3pm               141851 non-null float64
21  RainToday             142199 non-null object
22  RainTomorrow          142193 non-null object
dtypes: float64(16), object(7)
memory usage: 25.5+ MB

```

There are some missing values in the data set where both numeric and categorical value are missing

Visualisaition for cleaning

```

In [5]: # as our target value is raintomorrow will hapen r not
        # Lets check wheather its balance or not

```

```

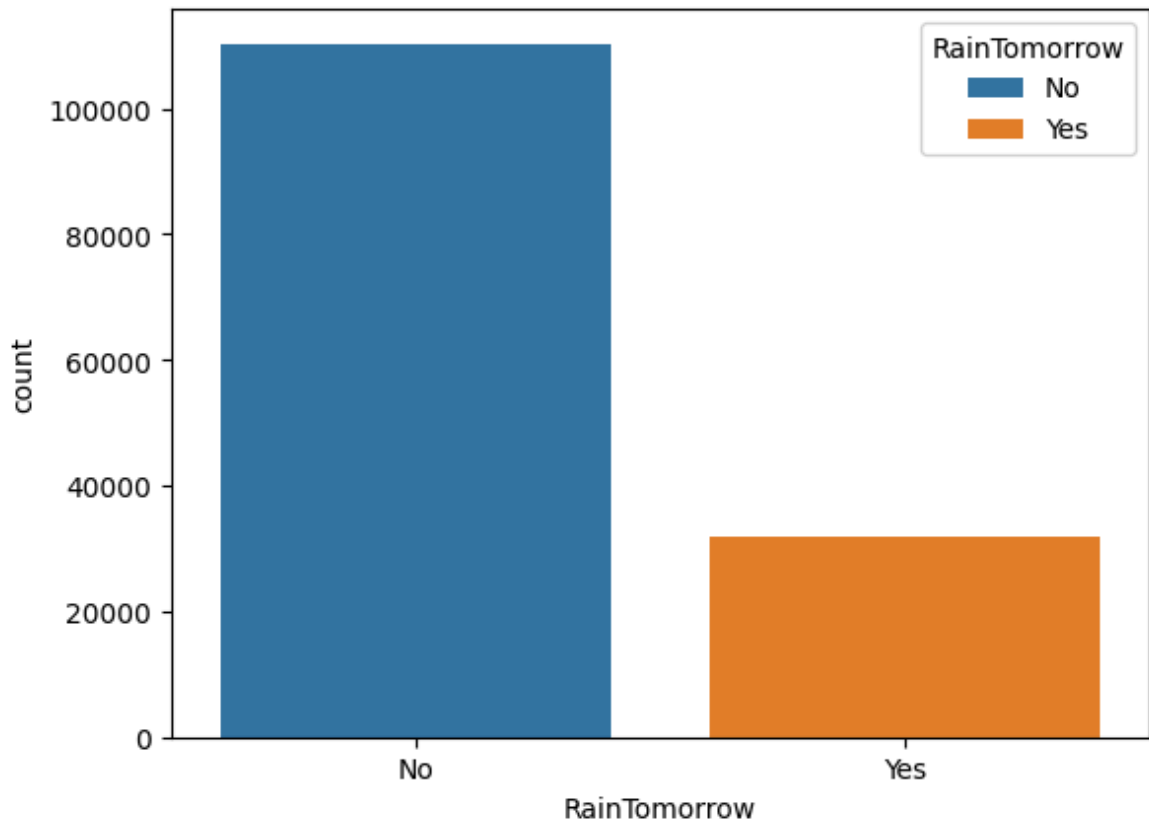
In [6]: sns.countplot(x=data['RainTomorrow'],hue=data["RainTomorrow"])

```

```


Out[6]: <Axes: xlabel='RainTomorrow', ylabel='count'>

```



```
In [7]: # take all numerical value one side to find the correlation between them
numeric_data=data[["MinTemp","MaxTemp","Rainfall","Evaporation","Sunshine","Wind
                 "Humidity3pm","Pressure9am","Pressure3pm","Cloud9am","Cloud3p
cor = numeric_data.corr() # find correlation
plt.subplots(figsize=(18,18)) # figure size
sns.heatmap(cor,annot=True,square=True) # represnted by heat map
```

Out[7]: <Axes: >



Encode date to months and day with sin and cos combination to get cyclic continuous feature

```
In [8]: # Parsing Date time
lengths = data["Date"].str.len()
lengths.value_counts()
```

```
Out[8]: Date
10      145460
Name: count, dtype: int64
```

```
In [9]: #There don't seem to be any error in dates so parsing values into datetime
data['Date'] = pd.to_datetime(data["Date"])
#Creating a column of year
data['year'] = data.Date.dt.year

# function to encode datetime into cyclic parameters.
#As I am planning to use this data in a neural network I prefer the months and d

def encode(data, col, max_val):
    data[col + '_sin'] = np.sin(2 * np.pi * data[col]/max_val)
    data[col + '_cos'] = np.cos(2 * np.pi * data[col]/max_val)
    return data
```

```
data['month'] = data.Date.dt.month
data = encode(data, 'month', 12)


data['day'] = data.Date.dt.day
data = encode(data, 'day', 31)

data.head()
```

Out[9]:


	Date	Location	MinTemp	MaxTemp	Rainfall	Evaporation	Sunshine	WindGustDir
0	2008-12-01	Albury	13.4	22.9	0.6	NaN	NaN	W
1	2008-12-02	Albury	7.4	25.1	0.0	NaN	NaN	WNW
2	2008-12-03	Albury	12.9	25.7	0.0	NaN	NaN	WSW
3	2008-12-04	Albury	9.2	28.0	0.0	NaN	NaN	NE
4	2008-12-05	Albury	17.5	32.3	1.0	NaN	NaN	W

5 rows × 30 columns



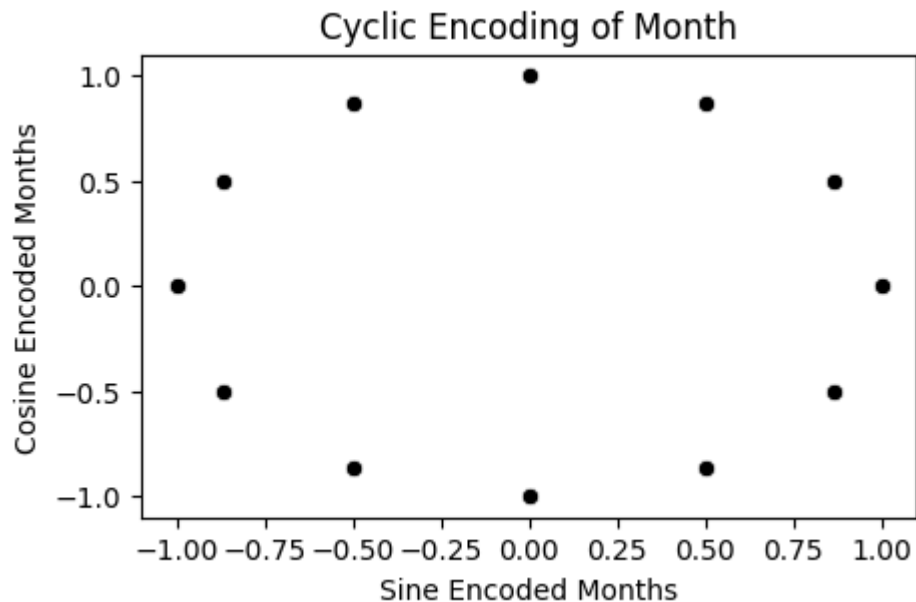
```
In [10]: # roughly a year's span section
section = data[:360]
plt.subplots(figsize=(5,3))
tm = section["day"].plot(color="red")
tm.set_title("Distribution Of Days Over Year")
tm.set_ylabel("Days In month")
tm.set_xlabel("Days In Year")
```

Out[10]: Text(0.5, 0, 'Days In Year')




```
In [11]: plt.subplots(figsize=(5,3))
cyclic_month = sns.scatterplot(x="month_sin",y="month_cos",data=data, color="black")
cyclic_month.set_title("Cyclic Encoding of Month")
cyclic_month.set_ylabel("Cosine Encoded Months")
cyclic_month.set_xlabel("Sine Encoded Months")
```

Out[11]: Text(0.5, 0, 'Sine Encoded Months')



```
In [12]: plt.subplots(figsize=(5,3))
cyclic_day = sns.scatterplot(x='day_sin',y='day_cos',data=data, color="black")
cyclic_day.set_title("Cyclic Encoding of Day")
cyclic_day.set_ylabel("Cosine Encoded Day")
cyclic_day.set_xlabel("Sine Encoded Day")
```

Out[12]: Text(0.5, 0, 'Sine Encoded Day')



Deal with missing value in the dataset

Categorical value missing treatments

```
In [13]: # Get list of categorical variables
s = (data.dtypes == "object")
object_cols = list(s[s].index)

print("Categorical variables:")
print(object_cols)
```

Categorical variables:
['Location', 'WindGustDir', 'WindDir9am', 'WindDir3pm', 'RainToday', 'RainTomorrow']

```
In [14]: # check Missing values in categorical variables

for i in object_cols:
    print(i, data[i].isnull().sum())
```

Location 0
WindGustDir 10326
WindDir9am 10566
WindDir3pm 4228
RainToday 3261
RainTomorrow 3267

```
In [15]: # Filling missing values with mode of the column in value as categorical

for i in object_cols:
    data[i].fillna(data[i].mode()[0])
```

Numeric value attribute in the dataste

```
In [16]: # Get list of neumeric variables
t = (data.dtypes == "float64")
num_cols = list(t[t].index)

print("Neumeric variables:")
print(num_cols)
```

Neumeric variables:
['MinTemp', 'MaxTemp', 'Rainfall', 'Evaporation', 'Sunshine', 'WindGustSpeed', 'WindSpeed9am', 'WindSpeed3pm', 'Humidity9am', 'Humidity3pm', 'Pressure9am', 'Pressure3pm', 'Cloud9am', 'Cloud3pm', 'Temp9am', 'Temp3pm', 'month_sin', 'month_cos', 'day_sin', 'day_cos']

```
In [17]: # checking Missing values in numeric value attribute

for i in num_cols:
    print(i, data[i].isnull().sum())
```


MinTemp 1485
MaxTemp 1261
Rainfall 3261
Evaporation 62790
Sunshine 69835
WindGustSpeed 10263
WindSpeed9am 1767
WindSpeed3pm 3062
Humidity9am 2654
Humidity3pm 4507
Pressure9am 15065
Pressure3pm 15028
Cloud9am 55888
Cloud3pm 59358
Temp9am 1767
Temp3pm 3609
month_sin 0
month_cos 0
day_sin 0
day_cos 0

In [18]: *# Filling missing values with median of the column in value*

```
for i in num_cols:  
    data[i].fillna(data[i].median())  
  
data.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 145460 entries, 0 to 145459
Data columns (total 30 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Date                  145460 non-null  datetime64[ns]
1   Location              145460 non-null  object
2   MinTemp               143975 non-null  float64
3   MaxTemp               144199 non-null  float64
4   Rainfall              142199 non-null  float64
5   Evaporation           82670 non-null   float64
6   Sunshine              75625 non-null   float64
7   WindGustDir           135134 non-null   object
8   WindGustSpeed         135197 non-null   float64
9   WindDir9am            134894 non-null   object
10  WindDir3pm            141232 non-null   object
11  WindSpeed9am          143693 non-null   float64
12  WindSpeed3pm          142398 non-null   float64
13  Humidity9am           142806 non-null   float64
14  Humidity3pm           140953 non-null   float64
15  Pressure9am           130395 non-null   float64
16  Pressure3pm           130432 non-null   float64
17  Cloud9am              89572 non-null    float64
18  Cloud3pm              86102 non-null    float64
19  Temp9am               143693 non-null   float64
20  Temp3pm               141851 non-null   float64
21  RainToday             142199 non-null   object
22  RainTomorrow          142193 non-null   object
23  year                  145460 non-null   int32
24  month                 145460 non-null   int32
25  month_sin             145460 non-null   float64
26  month_cos             145460 non-null   float64
27  day                   145460 non-null   int32
28  day_sin               145460 non-null   float64
29  day_cos               145460 non-null   float64
dtypes: datetime64[ns](1), float64(20), int32(3), object(6)
memory usage: 31.6+ MB

```

Successfully fill the missing value

visualisation

```


In [19]: # line plot to see out year wise rain fall
plt.subplots(figsize=(5,3))
sns.lineplot(x=data["year"],y=data["Rainfall"])

```

```


Out[19]: <Axes: xlabel='year', ylabel='Rainfall'>

```



```
In [20]: # wind speed by the year
plt.subplots(figsize=(10,5))
sns.barplot(x=data["year"],y=data["WindGustSpeed"],palette='pastel')
```

Out[20]: <Axes: xlabel='year', ylabel='WindGustSpeed'>



Data preprocessing

- Label encoder
- Scaling the data
- Detecting outlier
- Drop outlier

```
In [21]: # LabelEncoder
# Apply Label encoder to each column with categorical data
label_encoder = LabelEncoder()
for i in object_cols:
    data[i] = label_encoder.fit_transform(data[i])
```

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 145460 entries, 0 to 145459
Data columns (total 30 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Date                  145460 non-null  datetime64[ns]
1   Location              145460 non-null  int32
2   MinTemp               143975 non-null  float64
3   MaxTemp               144199 non-null  float64
4   Rainfall              142199 non-null  float64
5   Evaporation           82670 non-null   float64
6   Sunshine              75625 non-null   float64
7   WindGustDir           145460 non-null  int32
8   WindGustSpeed         135197 non-null  float64
9   WindDir9am            145460 non-null  int32
10  WindDir3pm            145460 non-null  int32
11  WindSpeed9am          143693 non-null  float64
12  WindSpeed3pm          142398 non-null  float64
13  Humidity9am           142806 non-null  float64
14  Humidity3pm           140953 non-null  float64
15  Pressure9am           130395 non-null  float64
16  Pressure3pm           130432 non-null  float64
17  Cloud9am              89572 non-null   float64
18  Cloud3pm              86102 non-null   float64
19  Temp9am               143693 non-null  float64
20  Temp3pm               141851 non-null  float64
21  RainToday             145460 non-null  int32
22  RainTomorrow          145460 non-null  int32
23  year                  145460 non-null  int32
24  month                 145460 non-null  int32
25  month_sin             145460 non-null  float64
26  month_cos             145460 non-null  float64
27  day                   145460 non-null  int32
28  day_sin               145460 non-null  float64
29  day_cos               145460 non-null  float64
dtypes: datetime64[ns](1), float64(20), int32(9)
memory usage: 28.3 MB
```

```
In [22]: # Scaling the data
# Preparing attributes of scale data

features = data.drop(['RainTomorrow', 'Date', 'day', 'month'], axis=1) # dropping

target = data['RainTomorrow']

#Set up a standard scaler for the features
col_names = list(features.columns)
s_scaler = StandardScaler()
features = s_scaler.fit_transform(features)
features = pd.DataFrame(features, columns=col_names)


features.describe().T
```

Out[22]:

	count	mean	std	min	25%	50%	7
Location	145460.0	7.815677e-18	1.000003	-1.672228	-0.899139	0.014511	0.857
MinTemp	143975.0	-3.790219e-16	1.000003	-3.234215	-0.717989	-0.030325	0.735
MaxTemp	144199.0	-1.387588e-16	1.000003	-3.936122	-0.747483	-0.087280	0.699
Rainfall	142199.0	6.875624e-17	1.000004	-0.278475	-0.278475	-0.278475	-0.184
Evaporation	82670.0	1.732738e-16	1.000006	-1.303922	-0.683942	-0.159343	0.460
Sunshine	75625.0	1.443165e-16	1.000007	-2.010636	-0.742625	0.208382	0.789
WindGustDir	145460.0	6.252542e-18	1.000003	-1.670768	-0.866215	0.139476	0.944
WindGustSpeed	135197.0	-1.731198e-16	1.000004	-2.501301	-0.664013	-0.076081	0.585
WindDir9am	145460.0	4.064152e-17	1.000003	-1.614034	-1.004491	0.011413	0.824
WindDir3pm	145460.0	-7.503050e-17	1.000003	-1.688306	-0.844398	-0.000489	0.843
WindSpeed9am	143693.0	-3.560304e-17	1.000003	-1.575197	-0.790034	-0.117037	0.555
WindSpeed3pm	142398.0	1.812309e-16	1.000004	-2.118405	-0.642770	0.038292	0.605
Humidity9am	142806.0	-2.292747e-16	1.000004	-3.619763	-0.624351	0.058814	0.741
Humidity3pm	140953.0	4.516727e-17	1.000004	-2.478339	-0.699136	0.022162	0.695
Pressure9am	130395.0	-1.277457e-14	1.000004	-5.227598	-0.668393	-0.007027	0.668
Pressure3pm	130432.0	-9.532011e-15	1.000004	-5.421883	-0.690013	-0.007942	0.674
Cloud9am	89572.0	1.167685e-16	1.000006	-1.540437	-1.194074	0.191379	0.884
Cloud3pm	86102.0	-2.376673e-17	1.000006	-1.657854	-0.922653	0.180150	0.915
Temp9am	143693.0	3.797658e-17	1.000003	-3.728099	-0.722889	-0.044790	0.710
Temp3pm	141851.0	5.770458e-16	1.000004	-3.904404	-0.732833	-0.084103	0.679
RainToday	145460.0	9.378812e-18	1.000003	-0.539860	-0.539860	-0.539860	-0.539

	count	mean	std	min	25%	50%	7
year	145460.0	2.080221e-14	1.000003	-2.273637	-0.697391	0.090732	0.878
month_sin	145460.0	5.861758e-19	1.000003	-1.434333	-0.725379	-0.016425	0.692
month_cos	145460.0	-2.745257e-17	1.000003	-1.388032	-1.198979	0.023080	0.728
day_sin	145460.0	1.075877e-17	1.000003	-1.403140	-1.019170	-0.003198	1.012
day_cos	145460.0	-1.353700e-17	1.000003	-1.392587	-1.055520	-0.044639	1.011

```
In [23]: # Detecting outliers
#looking at the scaled features
colours = ["red", "blue", "green", "yellow", "pink"]
plt.figure(figsize=(20,10))
sns.boxenplot(data = features,palette = colours)
plt.xticks(rotation=90)
plt.show()
```



```
In [24]: # Drop outlier
#full data for
features["RainTomorrow"] = target

#Dropping with outlier

features = features[(features["MinTemp"]<2.3)&(features["MinTemp"]>-2.3)]
features = features[(features["MaxTemp"]<2.3)&(features["MaxTemp"]>-2)]
features = features[(features["Rainfall"]<4.5)]
features = features[(features["Evaporation"]<2.8)]
features = features[(features["Sunshine"]<2.1)]
features = features[(features["WindGustSpeed"]<4)&(features["WindGustSpeed"]>-4)]
features = features[(features["WindSpeed9am"]<4)]
```

```

features = features[(features["WindSpeed3pm"]<2.5)]
features = features[(features["Humidity9am"]>-3)]
features = features[(features["Humidity3pm"]>-2.2)]
features = features[(features["Pressure9am"]< 2)&(features["Pressure9am"]>-2.7)]
features = features[(features["Pressure3pm"]< 2)&(features["Pressure3pm"]>-2.7)]
features = features[(features["Cloud9am"]<1.8)]
features = features[(features["Cloud3pm"]<2)]
features = features[(features["Temp9am"]<2.3)&(features["Temp9am"]>-2)]
features = features[(features["Temp3pm"]<2.3)&(features["Temp3pm"]>-2)]

features.shape

```


Out[24]: (52851, 27)

In [25]: features.head()

Out[25]:

	Location	MinTemp	MaxTemp	Rainfall	Evaporation	Sunshine	WindGustDir
6049	-0.96942	0.891770	1.682626	-0.278475	1.557527	1.238641	0.541753
6050	-0.96942	0.969914	0.797673	-0.278475	2.225199	1.423559	-0.061662
6052	-0.96942	1.126201	2.019750	-0.278475	1.271382	0.789554	-0.665077
6053	-0.96942	1.516919	2.132125	-0.278475	1.414455	1.212224	1.145167
6056	-0.96942	1.735721	1.514063	-0.278475	1.032928	1.317891	0.541753

5 rows × 27 columns




In [26]:

```

# visualise afer drop outlier
#looking at the scaled features without outliers

plt.figure(figsize=(20,10))
sns.boxenplot(data = features,palette = colours)
plt.xticks(rotation=90)
plt.show()

```



Model building

In this project, we build an artificial neural network.

Following steps are involved in the model building

- Assigning X and y the status of attributes and tags
- Splitting test and training sets
- Initialising the neural network
- Defining by adding layers
- Compiling the neural network
- Train the neural network

```
In [27]: # dependent and independent variable
X = features.drop(["RainTomorrow"], axis=1)
y = features["RainTomorrow"]

# Splitting test and training sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state=42)

X.shape
```

Out[27]: (52851, 26)

```
In [28]: # Initialising the ANN
model = Sequential()

# layers

model.add(Dense(units = 32, kernel_initializer = 'uniform', activation = 'relu',
```



```

model.add(Dense(units = 32, kernel_initializer = 'uniform', activation = 'relu'))
model.add(Dense(units = 16, kernel_initializer = 'uniform', activation = 'relu'))
model.add(Dropout(0.25))
model.add(Dense(units = 8, kernel_initializer = 'uniform', activation = 'relu'))
model.add(Dropout(0.5))
model.add(Dense(units = 1, kernel_initializer = 'uniform', activation = 'sigmoid'))

# Compiling the ANN
opt = Adam(learning_rate=0.00009)
model.compile(optimizer = opt, loss = 'binary_crossentropy', metrics = ['accuracy'])

# Train the ANN with 15 epochs
history = model.fit(X_train, y_train, batch_size = 32, epochs = 15, validation_s

```

```

Epoch 1/15
1057/1057 ————— 9s 4ms/step - accuracy: 0.7805 - loss: 0.6223 - val_accuracy: 0.7835 - val_loss: 0.3856
Epoch 2/15
1057/1057 ————— 4s 4ms/step - accuracy: 0.7817 - loss: 0.4125 - val_accuracy: 0.7835 - val_loss: 0.3747
Epoch 3/15
1057/1057 ————— 4s 4ms/step - accuracy: 0.7993 - loss: 0.4039 - val_accuracy: 0.8479 - val_loss: 0.3670
Epoch 4/15
1057/1057 ————— 5s 4ms/step - accuracy: 0.8142 - loss: 0.3956 - val_accuracy: 0.8491 - val_loss: 0.3615
Epoch 5/15
1057/1057 ————— 5s 5ms/step - accuracy: 0.8127 - loss: 0.3921 - val_accuracy: 0.8499 - val_loss: 0.3564
Epoch 6/15
1057/1057 ————— 4s 4ms/step - accuracy: 0.8104 - loss: 0.3910 - val_accuracy: 0.8512 - val_loss: 0.3507
Epoch 7/15
1057/1057 ————— 5s 4ms/step - accuracy: 0.8142 - loss: 0.3902 - val_accuracy: 0.8521 - val_loss: 0.3474
Epoch 8/15
1057/1057 ————— 5s 4ms/step - accuracy: 0.8191 - loss: 0.3759 - val_accuracy: 0.8530 - val_loss: 0.3462
Epoch 9/15
1057/1057 ————— 5s 4ms/step - accuracy: 0.8106 - loss: 0.3873 - val_accuracy: 0.8550 - val_loss: 0.3432
Epoch 10/15
1057/1057 ————— 5s 5ms/step - accuracy: 0.8156 - loss: 0.3859 - val_accuracy: 0.8551 - val_loss: 0.3419
Epoch 11/15
1057/1057 ————— 3s 3ms/step - accuracy: 0.8096 - loss: 0.3909 - val_accuracy: 0.8561 - val_loss: 0.3411
Epoch 12/15
1057/1057 ————— 6s 4ms/step - accuracy: 0.8231 - loss: 0.3752 - val_accuracy: 0.8569 - val_loss: 0.3414
Epoch 13/15
1057/1057 ————— 4s 4ms/step - accuracy: 0.8199 - loss: 0.3794 - val_accuracy: 0.8556 - val_loss: 0.3388
Epoch 14/15
1057/1057 ————— 6s 5ms/step - accuracy: 0.8159 - loss: 0.3847 - val_accuracy: 0.8566 - val_loss: 0.3389
Epoch 15/15
1057/1057 ————— 5s 5ms/step - accuracy: 0.8215 - loss: 0.3799 - val_accuracy: 0.8567 - val_loss: 0.3379


```

Plotting training and validation loss over epochs

```
In [29]: history_df = pd.DataFrame(history.history)

plt.plot(history_df.loc[:, ['loss']], "red", label='Training loss')
plt.plot(history_df.loc[:, ['val_loss']], "blue", label='Validation loss')
plt.title('Training and Validation loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend(loc="best")

plt.show()
```




Plotting training and validation accuracy over epochs

```
In [30]: history_df = pd.DataFrame(history.history)

plt.plot(history_df.loc[:, ['accuracy']], "yellow", label='Training accuracy')
plt.plot(history_df.loc[:, ['val_accuracy']], "blue", label='Validation accuracy')

plt.title('Training and Validation accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.show()
```



Conclusion


- Testing on the test set
- Evaluating the confusion matrix
- Evaluating the classification report

```
In [31]: # Predicting the test set results
y_pred = model.predict(X_test)
y_pred = (y_pred > 0.5)
```

331/331 ————— 1s 4ms/step

```
In [32]: # confusion matrix
cmap1 = sns.diverging_palette(260,-10,s=50, l=75, n=5, as_cmap=True)
plt.subplots(figsize=(8,5))
cf_matrix = confusion_matrix(y_test, y_pred)
sns.heatmap(cf_matrix/np.sum(cf_matrix), cmap = cmap1, annot = True, annot_kws =
```

Out[32]: <Axes: >



```
In [33]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.88	0.95	0.91	8296
1	0.75	0.51	0.60	2254
2	0.00	0.00	0.00	21
accuracy			0.86	10571
macro avg	0.54	0.49	0.51	10571
weighted avg	0.85	0.86	0.84	10571

The model train with 85% accuracy

complete

```
In [ ]:
```